

REINFORCEMENT LEARNING

Seminar @ UPC TelecomBCN Barcelona (2nd edition). Spring 2020.



Instructors



Josep
Vidal



Margarita
Cabrera



Xavier
Giró-i-Nieto

Organizers



UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH



Supporters



Google Cloud



GitHub Education

+ info: <http://bit.ly/upcrl-2020>

<https://telecombcn-dl.github.io/mrl-2020/>

MRL 2020 - Day 9 - Part 1

How to train your neural network



Xavier Giro-i-Nieto



@DocXavi



xavier.giro@upc.edu

Associate Professor

ETSETB TelecomBCN

Universitat Politècnica de Catalunya



Acknowledgments



Víctor Campos

victor.campos@bsc.es

PhD Candidate

Barcelona Supercomputing Center



Míriam Bellver

miriam.bellver@bsc.edu

PhD Candidate

Barcelona Supercomputing Center



Kevin McGuinness

kevin.mcguinness@dcu.ie

Research Fellow

Insight Centre for Data Analytics
Dublin City University



Elisa Sayrol

elisa.sayrol@upc.edu

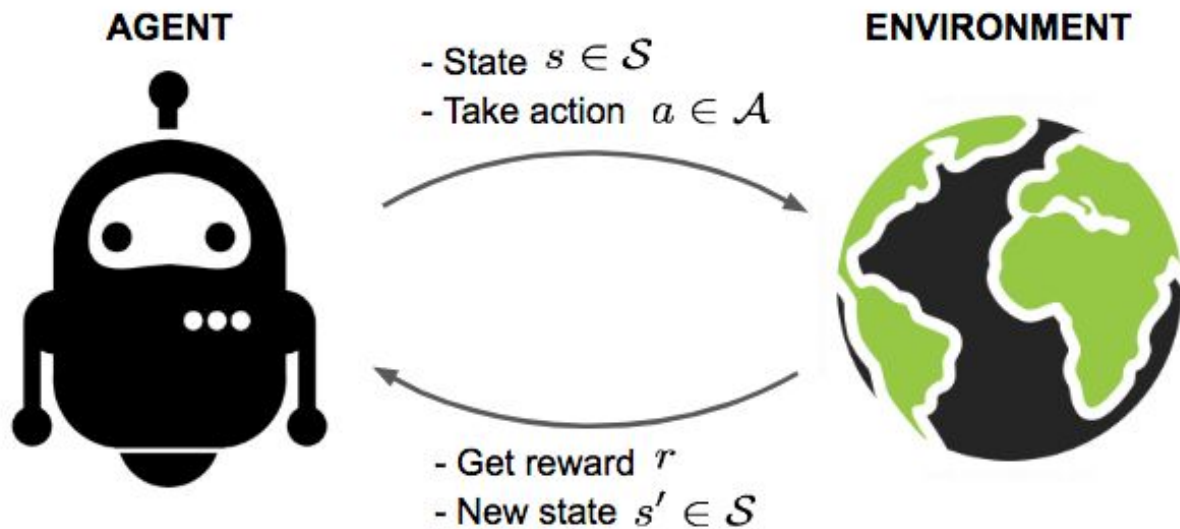
Associate Professor

ETSETB TelecomBCN
Universitat Politècnica de Catalunya

Outline

1. RL with Neural Networks
2. Loss functions
3. Backpropagation
4. Optimizers

Reinforcement Learning (with extrinsic reward)



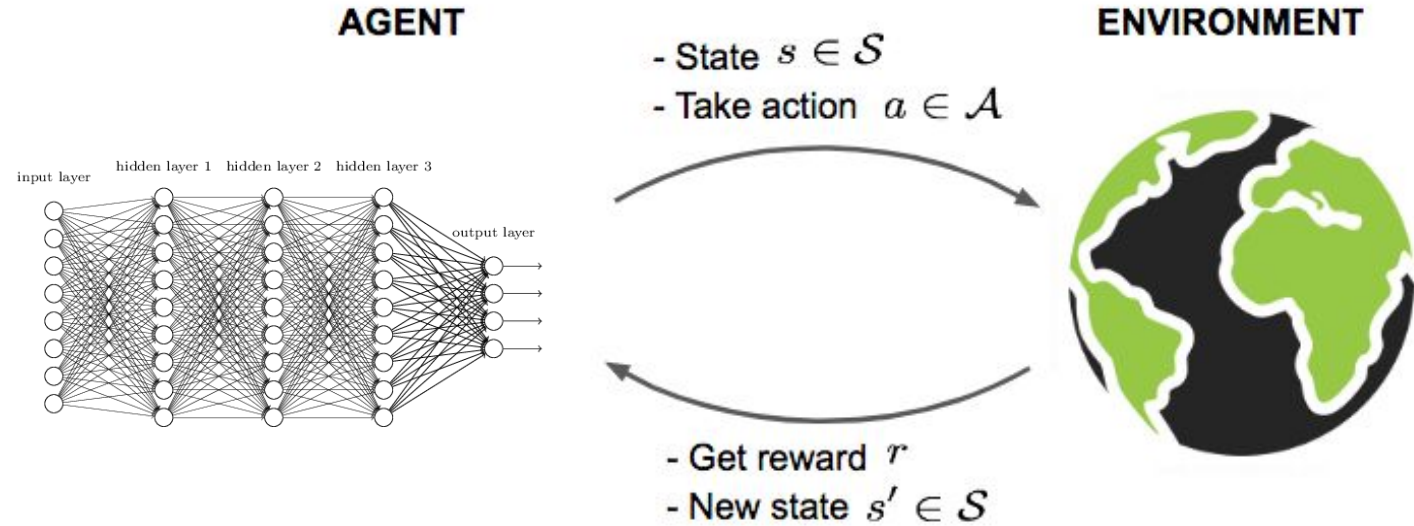
Policy π

Value
function

Model
(of the Environment)

Goals of Reinforcement Learning

Reinforcement Learning with Neural Networks (NN)



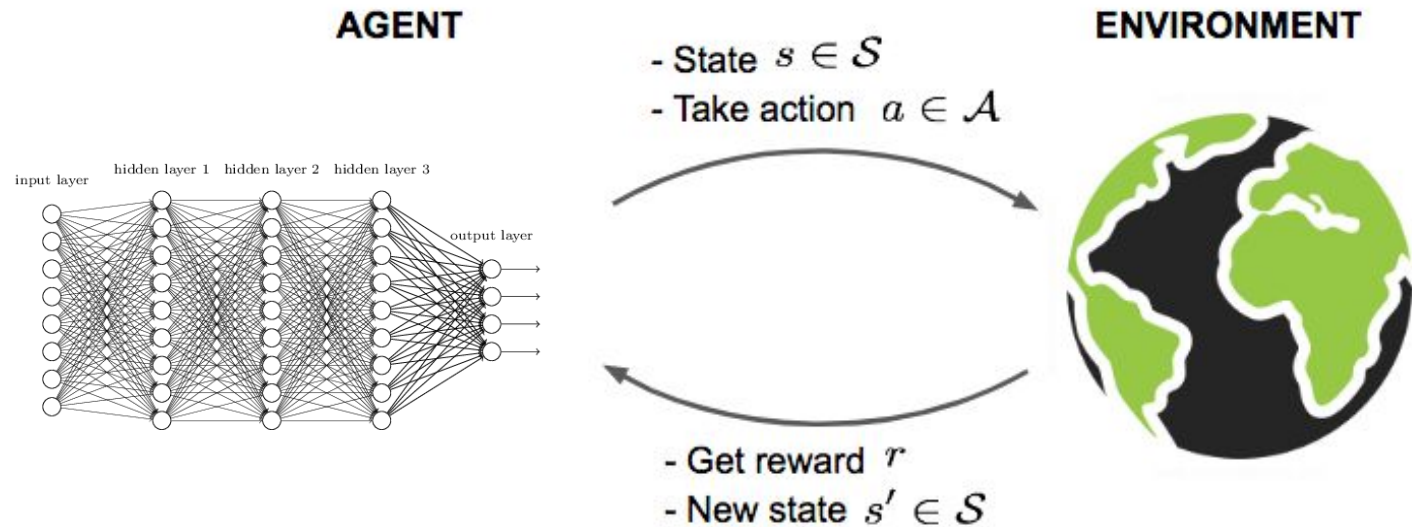
Policy π

Value
function

Model
(of the Environment)

Goals of Reinforcement Learning

Policy-based RL with Neural Networks (NN)



Policy π

Value
function


Model
(of the Environment)

Goals of Reinforcement Learning

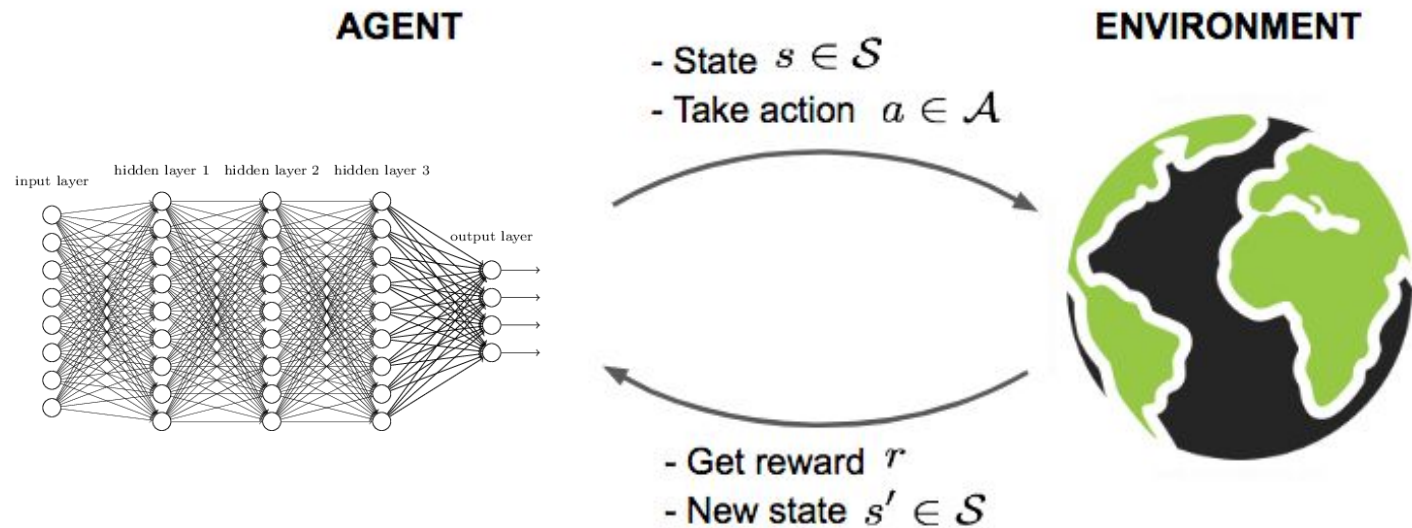
Policy-based RL with Neural Networks (NN)

The Policy π is a function $S \rightarrow A$ that specifies which action to take in each state:

A Multi-Layer Perceptron (MLP) can implement a classifier to predict the distribution of actions given a state.



Value-based RL with Neural Networks (NN)



Policy π

Value
function


Model
(of the Environment)

Goals of Reinforcement Learning

Value-based RL with Neural Networks (NN)


The action-value $Q_{\pi}(s,a)$ function is the expected return for taking action a when being in state s .

A Multi-Layer Perceptron (MLP) can implement a regressor to predict $Q_{\pi}(s,a)$.




Single neuron model (perceptron)

Which components of the neurons must be estimated during training ?




Single neuron model (perceptron)

The **weights** (w_i) and **bias** (b) must be estimated during training.



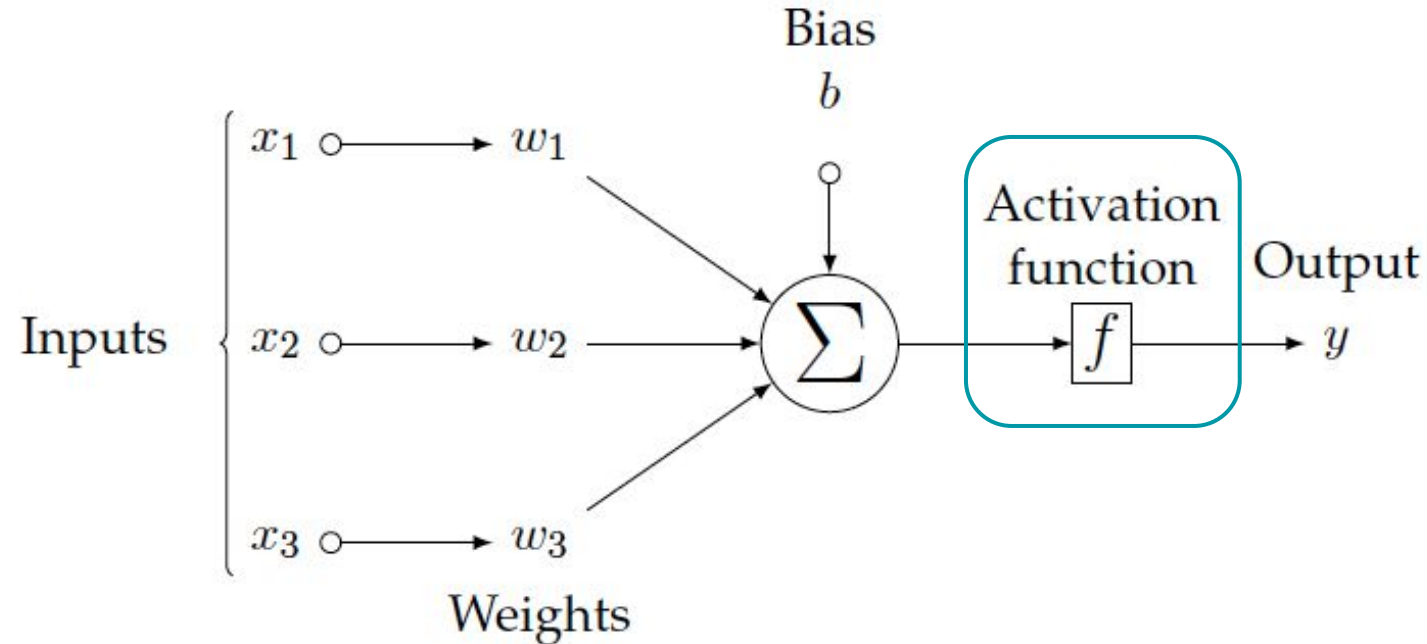
Single neuron model (perceptron)

The **activation function (f)** is a design choice.



Single neuron model (perceptron)

The **activation function (f)** is a design choice.



Single neuron model (perceptron)

Activation functions:


- They act as a **threshold**

Desirable properties

- Mostly smooth, continuous, differentiable
- Fairly linear


Common nonlinearities

- Sigmoid
- Tanh
- ReLU = $\max(0, x)$







Why do we need them?

If we only use linear layers we are only able to learn linear transformations of our input.



Single neuron model (perceptron)

Other popular activation functions:



Multi-Layer Perceptron (MLP)

When each node in each layer is a linear combination of **all inputs from the previous layer** then the network is called a multilayer perceptron (MLP)


Weights can be organized into matrices.

Forward pass computes

$$\mathbf{h}_0 = \mathbf{x}$$

$$\mathbf{h}^{(t)} = g(W^{(t)}\mathbf{h}^{(t-1)} + \mathbf{b}^{(t)})$$

$$f(\mathbf{x}) = \mathbf{h}^{(L)}$$



Multi-Layer Perceptron (MLP)

W_1				h_0		b_1	
w_{11}	w_{12}	w_{13}	w_{14}	x_1		b_1	
w_{21}	w_{22}	w_{23}	w_{24}	x_2		b_2	
w_{31}	w_{32}	w_{33}	w_{34}	x_3		b_3	
w_{41}	w_{42}	w_{43}	w_{44}	x_4		b_4	


Forward pass computes

$$h_0 = x$$

$$h^{(t)} = g(W^{(t)}h^{(t-1)} + b^{(t)})$$

$$f(x) = h^{(L)}$$

$$h_{11} = g(\mathbf{w}\mathbf{x} + \mathbf{b})$$



Multi-Layer Perceptron (MLP)

W_1				h_0		b_1	
w_{11}	w_{12}	w_{13}	w_{14}	x_1		b_1	
w_{21}	w_{22}	w_{23}	w_{24}	x_2		b_2	
w_{31}	w_{32}	w_{33}	w_{34}	x_3		b_3	
w_{41}	w_{42}	w_{43}	w_{44}	x_4		b_4	

Forward pass computes


$$h_0 = x$$

$$h^{(t)} = g(W^{(t)}h^{(t-1)} + b^{(t)})$$

$$f(x) = h^{(L)}$$

$$h_{11} = g(\mathbf{w}\mathbf{x} + \mathbf{b})$$

$$h_{12} = g(\mathbf{w}\mathbf{x} + \mathbf{b})$$




Multi-Layer Perceptron (MLP)

How many parameters contains the following MLP ?

Model


- 3 layer neural network (2 hidden layers)
- Tanh units (activation function)
- 512-512-10



Multi-Layer Perceptron (MLP)

How many parameters contains the following MLP ?

Layer	#Weights	#Biases	Total
1	784 x 512	512	401,920
2	512 x 512	512	262,656
3	512 x 10	10	5,130
			669,706



Multi-Layer Perceptron (MLP)

DEEP LEARNING FOR ARTIFICIAL INTELLIGENCE

Master Course UPC ETSETB TelecomBCN Barcelona, Autumn 2017



Instructors



Organizers



Supporters



+ info: <http://dlai.deeplearning.barcelona>

[\[course site\]](#)

Day 2 Lecture 1

Multilayer Perceptron



Elisa Sayrol




UNIVERSITAT POLITÈCNICA DE CATALUNYA
BARCELONATECH
Department of Signal Theory
and Communications
Image Processing Group

Outline

1. RL with Neural Networks
2. Loss functions
3. Backpropagation
4. Optimizers

Black box abstraction of supervised learning



Loss function - $L(y, \hat{y})$

loss function


= cost function

= objective function

= error function

Loss function - $L(y, \hat{y})$

How good does
our network with
the training data?

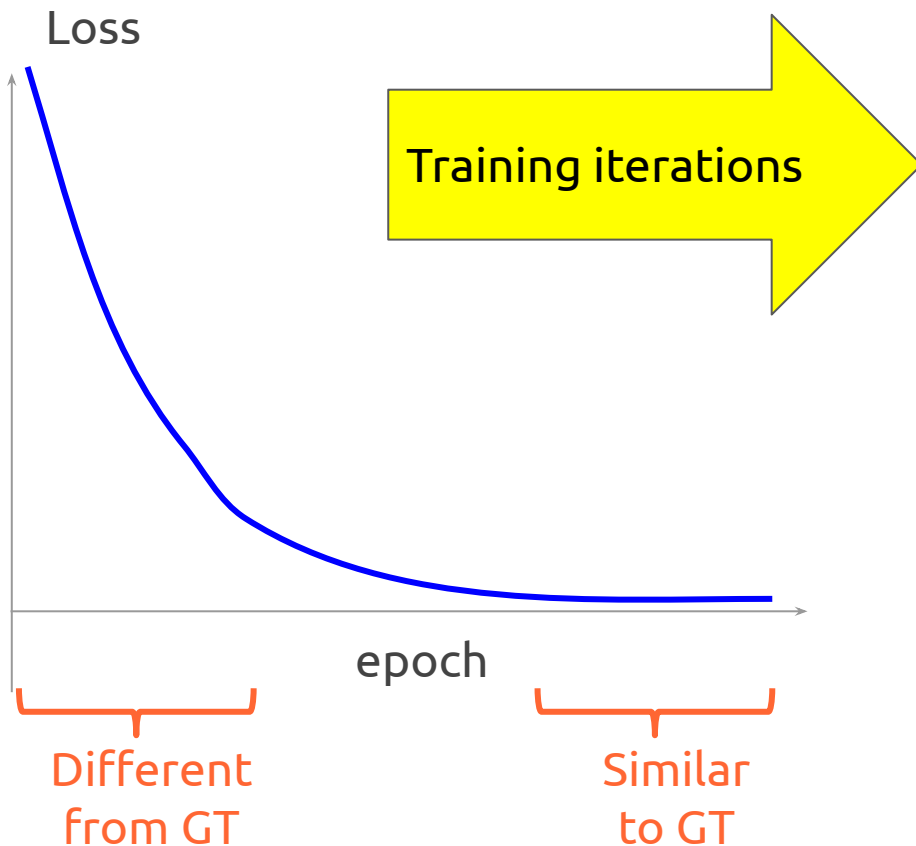


$$\mathcal{L}(w) = \text{distance}(f_{\theta}(x), y)$$

Labels (ground truth) y
input x
error $\mathcal{L}(w)$
parameters (weights, biases) θ

Loss function - $L(y, \hat{y})$

The Loss value should decrease when the more the NN output matches the ground truth (GT).



Loss functions for regression

Regression: the network predicts **continuous, numeric** variables

- Example: Price of a house

$$y = \mathbf{w}^T \cdot \mathbf{x} + b = w_1 \cdot x_1 + w_2 \cdot x_2 + w_3 \cdot x_3 + \dots + w_M \cdot x_M + b$$

Which loss function would you choose for this problem ?




Loss function - $L(y, \hat{y})$

The **loss function** assesses the performance of our model by comparing its predictions (\hat{y}) to an expected value (y), typically coming from annotations.

Example: the predicted price (\hat{y}) and one actually paid (y) could be compared with the Euclidean distance (also referred as L2 distance or Mean Square Error - MSE):

$$y = w_1 \cdot x_1 + w_2 \cdot x_2 + w_3 \cdot x_3 + b = \mathbf{w}^T \cdot \mathbf{x} + b$$




Loss functions for regression

L1 Loss

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n |y_i - f_{\theta}(x_i)|$$

L2 Loss

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (y_i - f_{\theta}(x_i))^2$$



Which of these two losses is lighter to compute ?


Loss functions for regression

L1 Loss

$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n |y_i - f_{\theta}(x_i)|$$

L2 Loss


$$\mathcal{L} = \frac{1}{n} \sum_{i=1}^n (y_i - f_{\theta}(x_i))^2$$



Which of these two losses is more sensitive to outliers ($|y_i - f_{\theta}(x_i)| \gg 1$) ?



Loss functions for regression

L2 loss has problems when there are outliers in the data



Loss functions for regression

Pseudo-Huber loss (Charbonnier, Smooth L_1) is generally more resilient to outliers than L2.



Outline

1. RL with Neural Networks
2. Loss functions
- 3. Backpropagation**
4. Optimizers


Backpropagation

Discussion: Consider the single parameter model...

$$\hat{y} = x \cdot w$$

...and that, given a pair (y, \hat{y}) , we would like to update the current w_t value to a new w_{t+1} based on the loss function L_w .

- (a) Would you increase or decrease w_t ?
- (b) What operation could indicate which way to go?
- (c) How much would you increase or decrease w_t ?



Gradient Descent (GD)


Motivation for this lecture:

if we had a way to estimate the gradient of the loss (∇L) with respect to the parameter(s), we could use gradient descent to optimize them.

$$w_{t+1} \leftarrow w_t - \alpha \nabla \mathcal{L}_w(w_t)$$

↓
Descend
(minus sign)

↑
Learning
rate (LR)



Gradient Descent (GD)

Backpropagation will allow us to compute the gradients of the loss function with respect to:

- all model parameters (**w** & **b**) - final goal during training
- input/intermediate data - visualization & interpretability purposes.

Gradients will **“flow”** from the output of the model towards the input (“back”).

Let the Gradient Flo

Celebrate NIPS 2017 with Intel AI

Join us for an exclusive party – and a surprise reveal.

Giveaways, buskers, acrobats, DJ Nostalgia B and a special performance by Flo Rida!

When


Tuesday, December 5th
9:00 PM – 12:00 AM
Door open at 9:00 PM
Show up early - space is limited

Where

The Loft on Pine
230 Pine Avenue
Long Beach, CA 90802
Near the Long Beach Convention Center




Computational graph of a simple perceptron





Question: What is the order of operations of this perceptron with a sigmoid activation?


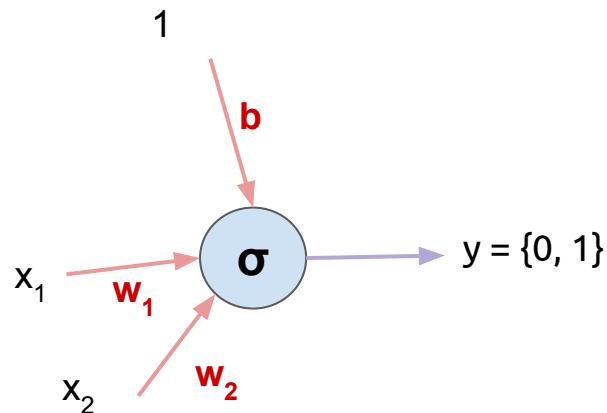
Computational graph of a perceptron





Computational graph of a perceptron





Computational graph of a perceptron



Computational graph of a perceptron



Computational graph of a perceptron




Computational graph of a perceptron

Challenge: How to compute the gradient of the loss function with respect to w_1 , w_2 or b ?

$$\frac{\partial \mathcal{L}(y, \hat{y})}{\partial w_1} = ?$$


$$\frac{\partial \mathcal{L}(y, \hat{y})}{\partial w_2} = ?$$

$$\frac{\partial \mathcal{L}(y, \hat{y})}{\partial b} = ?$$



Gradients from composition (chain rule)

$$\hat{y} = g_4(g_3(g_2(g_1(x_1))))$$



Decompose into steps (**forward propagation**):

$$x_2 = g_1(x_1)$$

$$x_3 = g_2(x_2)$$

$$x_4 = g_3(x_3)$$


$$\hat{y} = x_5 = g_4(x_4)$$



Forward pass

Gradients from composition (chain rule)

$$\hat{y} = g_4(g_3(g_2(g_1(x_1))))$$




Want to find $\frac{\partial \hat{y}}{\partial x_1}$. Chain rule:

$$\frac{\partial \hat{y}}{\partial x_1} = \frac{\partial \hat{y}}{\partial x_4} \frac{\partial x_4}{\partial x_3} \frac{\partial x_3}{\partial x_2} \frac{\partial x_2}{\partial x_1}$$

How does a variation
("difference") on the
input affect the
prediction?

Backward pass

Gradients from composition (chain rule)




Decompose into steps again. Let $\delta_k = \frac{\partial \hat{y}}{\partial x_k}$. **Backpropagation:**

$$\delta_5 = \frac{\partial \hat{y}}{\partial x_5} = 1$$

A variation in x_5
directly affects on \hat{y}
with a 1:1 factor.

Gradients from composition (chain rule)




Decompose into steps again. Let $\delta_k = \frac{\partial \hat{y}}{\partial x_k}$. **Backpropagation:**

$$\delta_5 = \frac{\partial \hat{y}}{\partial x_5} = 1$$

$$\delta_4 = \frac{\partial \hat{y}}{\partial x_4} = \frac{\partial \hat{y}}{\partial x_5} \frac{\partial x_5}{\partial x_4} = \delta_5 g'_4(x_4)$$

How does a variation on x_4 affect the predicted \hat{y} ?

Gradients from composition (chain rule)



Decompose into steps again. Let $\delta_k = \frac{\partial \hat{y}}{\partial x_k}$. **Backpropagation:**


$$\delta_5 = \frac{\partial \hat{y}}{\partial x_5} = 1$$

$$\delta_4 = \frac{\partial \hat{y}}{\partial x_4} = \frac{\partial \hat{y}}{\partial x_5} \frac{\partial x_5}{\partial x_4} = \delta_5 g'_4(x_4)$$

How does a variation on x_4 affect the predicted \hat{y} ?

It corresponds to how a variation of x_5 affects \hat{y} ...

Gradients from composition (chain rule)



Decompose into steps again. Let $\delta_k = \frac{\partial \hat{y}}{\partial x_k}$. **Backpropagation:**

$$\delta_5 = \frac{\partial \hat{y}}{\partial x_5} = 1$$


$$\delta_4 = \frac{\partial \hat{y}}{\partial x_4} = \frac{\partial \hat{y}}{\partial x_5} \frac{\partial x_5}{\partial x_4} = \delta_5 g'_4(x_4)$$

How does a variation on x_4 affect the predicted \hat{y} ?

It corresponds to how a variation of x_5 affects \hat{y} ...

...**multiplied** by how a variation near the input x_4 affects the output $g_4(x_4)$.

Gradients from composition (chain rule)



The same reasoning can be iteratively applied until reaching $\frac{\partial \hat{y}}{\partial x_1}$:

$$\delta_5 = \frac{\partial \hat{y}}{\partial x_5} = 1$$

$$\delta_4 = \frac{\partial \hat{y}}{\partial x_4} = \frac{\partial \hat{y}}{\partial x_5} \frac{\partial x_5}{\partial x_4} = \delta_5 g'_4(x_4)$$

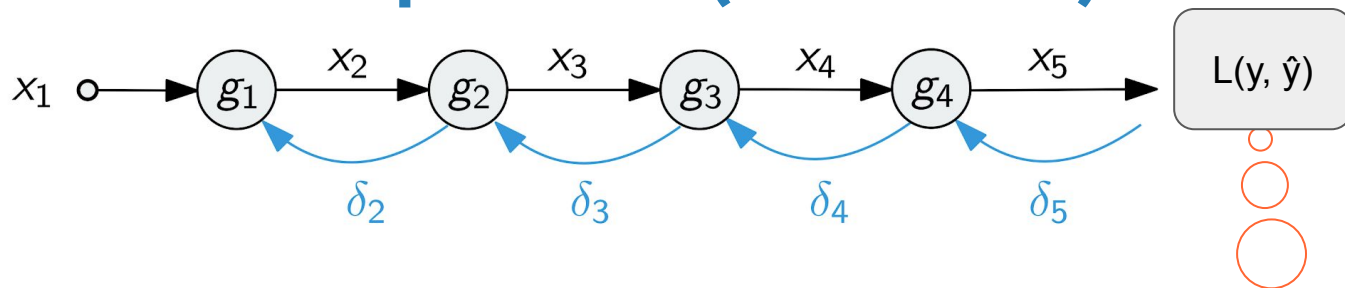
$$\delta_3 = \frac{\partial \hat{y}}{\partial x_3} = \frac{\partial \hat{y}}{\partial x_4} \frac{\partial x_4}{\partial x_3} = \delta_4 g'_3(x_3)$$

$$\delta_2 = \frac{\partial \hat{y}}{\partial x_2} = \frac{\partial \hat{y}}{\partial x_3} \frac{\partial x_3}{\partial x_2} = \delta_3 g'_2(x_2)$$

$$\delta_1 = \frac{\partial \hat{y}}{\partial x_1} = \frac{\partial \hat{y}}{\partial x_2} \frac{\partial x_2}{\partial x_1} = \delta_2 g'_1(x_1)$$

Backward pass

Gradients from composition (chain rule)




When training NN, we will actually compute the derivative over the loss function with respect the weights and biases

Backward pass

Gradients from composition (chain rule)

Question: What are the derivatives of the function involved in the computational graph of a perceptron ?

- SUM (+) $\frac{\partial(a + b)}{\partial a}$
- PRODUCT (x) $\frac{\partial(a \cdot b)}{\partial a}$
- SIGMOID (σ) $\frac{\partial \sigma(x)}{\partial x}$



Gradient weights for sigmoid σ

(*)

$$\frac{\partial \sigma(x)}{\partial x} = \frac{\partial}{\partial x} \left(\frac{1}{1 + e^{-x}} \right) = \frac{-1}{(1 + e^{-x})^2} \frac{\partial(1 + e^{-x})}{\partial x} = \frac{-1}{(1 + e^{-x})^2} \frac{\partial(e^{-x})}{\partial x}$$

$$\frac{\partial \sigma(x)}{\partial x} = \frac{-1}{(1 + e^{-x})^2} (-e^{-x}) = \frac{e^{-x}}{(1 + e^{-x})^2}$$




...which can be re-arranged as...

$$\frac{\partial \sigma(x)}{\partial x} = \frac{e^{-x}}{(1 + e^{-x})^2} = \frac{e^{-x}}{(1 + e^{-x})} \frac{1}{(1 + e^{-x})}$$

$$\frac{\partial \sigma(x)}{\partial x} = \left(\frac{1 + e^{-x}}{1 + e^{-x}} - \frac{1}{1 + e^{-x}} \right) \sigma(x)$$

$$\frac{\partial \sigma(x)}{\partial x} = (1 - \sigma(x)) \sigma(x)$$

(*) $f(x) = \frac{g(x)}{h(x)} \quad f'(x) = \frac{g'(x)h(x) - g(x)h'(x)}{h^2}$



Gradient backpropagation in a perceptron

$$\frac{\partial \sigma(x)}{\partial x} = (1 - \sigma(x)) \sigma(x)$$

```
import math


dot=1

# sigmoid function
f = 1.0 / (1 + math.exp(-dot))

# gradient on dot variable,
ddot = (1 - f) * f

print(ddot)
```

0.19661193324148185

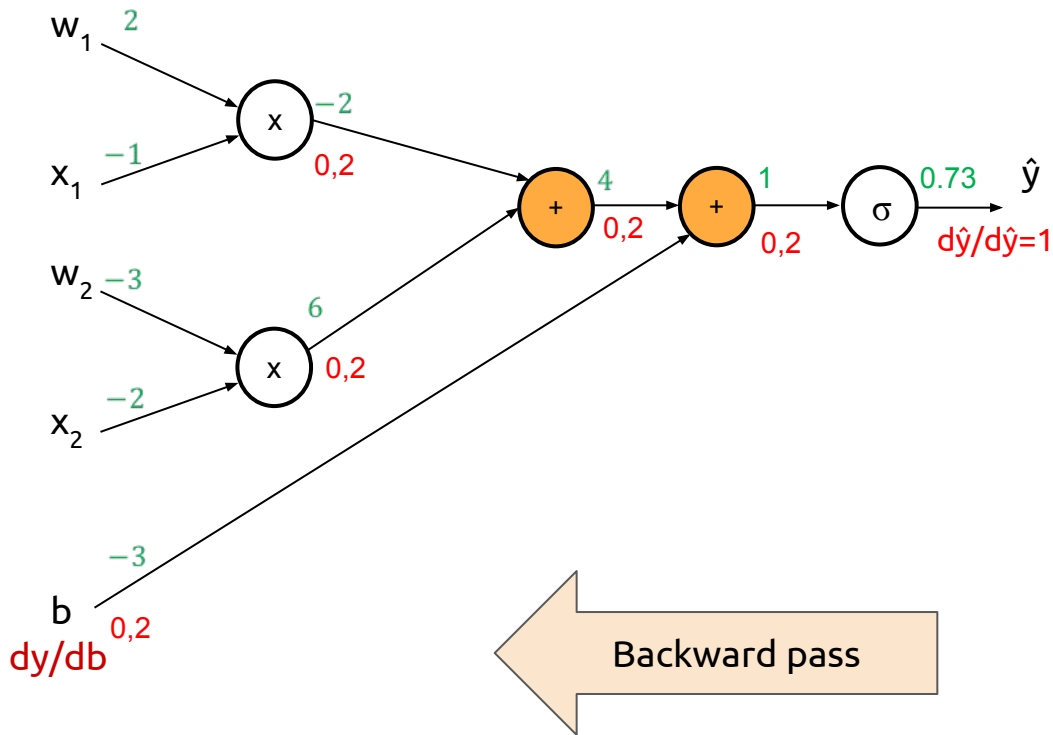


Gradient backpropagation in a perceptron

SUM

$$\frac{\partial(a + b)}{\partial a} = 1$$

$$\frac{\partial(a + b)}{\partial b} = 1$$

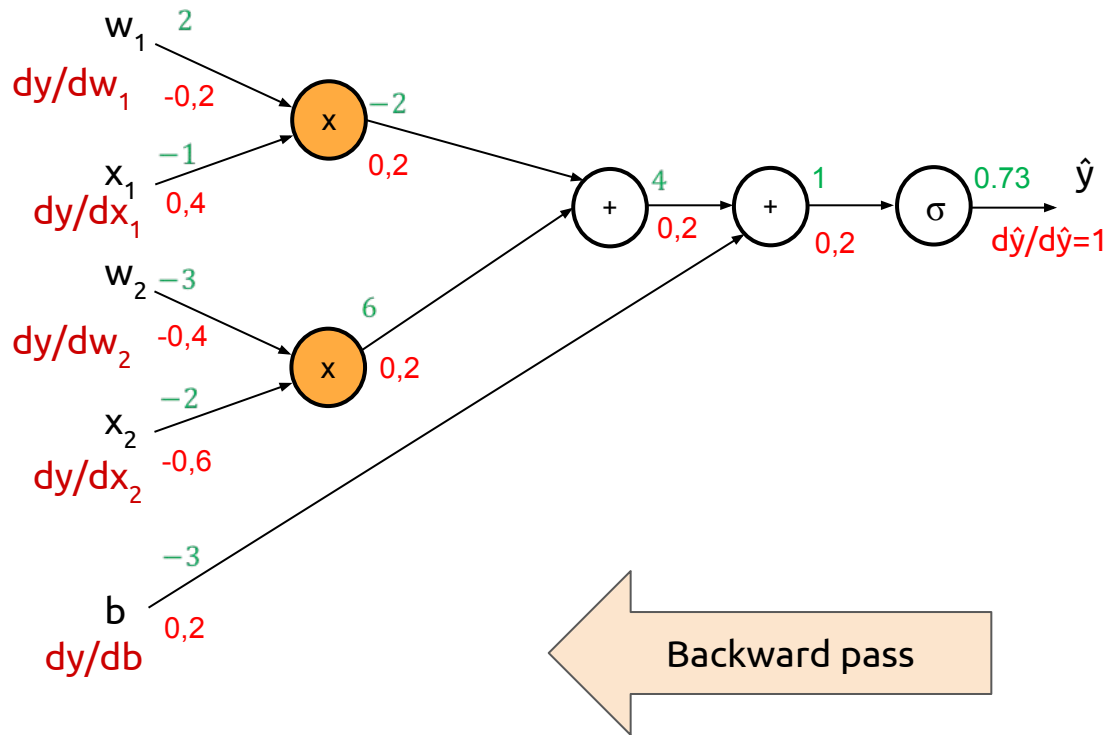


Gradient backpropagation in a perceptron

Product

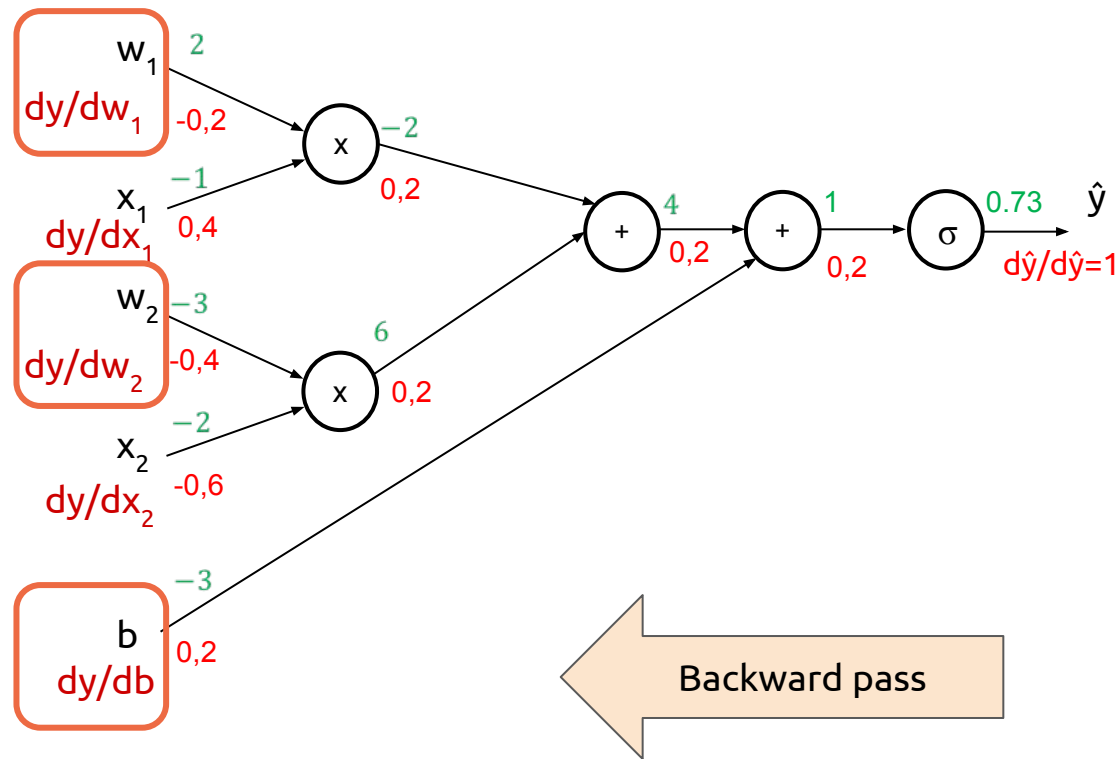
$$\frac{\partial(a \cdot b)}{\partial a} = b$$

$$\frac{\partial(a \cdot b)}{\partial b} = a$$




Gradient backpropagation in a perceptron

Normally, we will be interested only on the weights (w_i) and biases (b), not the inputs (x_i). The weights are the parameters to learn in our models.




(bonus) Gradients weights for MAX & SPLIT

Max: Routes the gradient only to the higher input branch (not sensitive to the lower branches).





Split: Branches that split in the forward pass and merge in the backward pass, add gradients



(bonus) Gradient weights for ReLU


$$\text{ReLU}(x) = \begin{cases} x & \text{if } x \geq 0 \\ 0 & \text{if } x < 0 \end{cases}$$

$$\frac{\partial \text{ReLU}(x)}{\partial x} = u(x) = \begin{cases} 1 & \text{if } x > 0 \\ 0 & \text{if } x \leq 0 \end{cases}$$



Backpropagation across layers

Gradients can flow across stacked layers of neurons to estimate their parameters.



Backpropagation & RL



Greg Brockman ✓

@gdb

Seguint

For differentiable problems, there's backpropagation. For everything else, there's RL.

Tradueix el tuit

18:11 - 31 de gen. de 2019



Yann LeCun

@ylecun

Seguint

En resposta a @gdb

Not quite right.

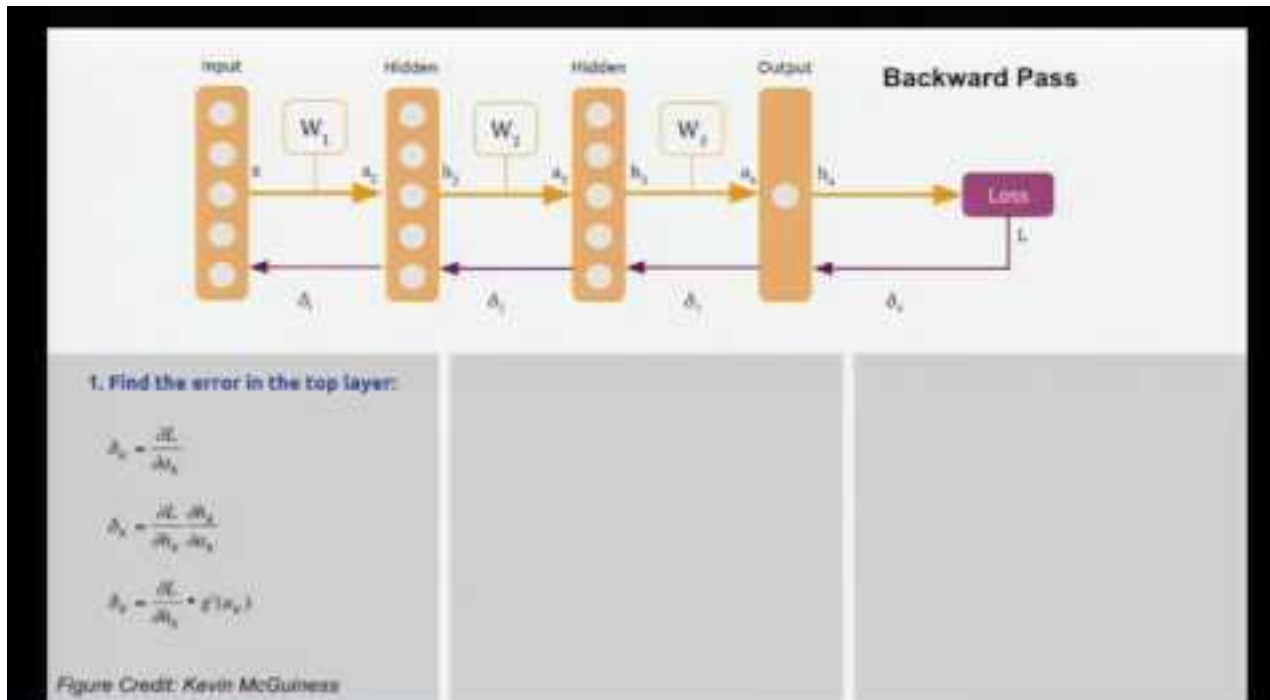
A more accurate statement would be "for everything else, there is gradient-free (zeroth-order) optimization."

RL is when there is a sequential decision process and what you see depends on previous actions you took.


Tradueix el tuit

2:38 - 1 de febr. de 2019

Backpropagation: Learn more



Backpropagation: Learn more



Gilbert Strang, [“27. Backpropagation: Find Partial Derivatives”](#). MIT 18.065 (2018)



Creative Commons, [“Yoshua Bengio Extra Footage 1: Brainstorm with students”](#) (2018)

Outline

1. RL with Neural Networks
2. Loss functions
3. Backpropagation
- 4. Optimizers**

Gradient Descent (GD)


Motivation for this lecture:

if we had a way to estimate the gradient of the loss (∇L) with respect to the parameter(s), we could use gradient descent to optimize them.

$$w_{t+1} \leftarrow w_t - \alpha \nabla \mathcal{L}_w(w_t)$$

↓
Descend
(minus sign)

↑
Learning
rate (LR)



Gradient descent (GD)

Computing the gradient for the full dataset at each step is slow

- Especially if the dataset is large!

For most losses we care about, the total loss can be expressed as a sum (or average) of losses on the individual examples

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N L(\mathbf{y}_i, \hat{\mathbf{y}}_i)$$

The gradient is the average of the gradients on individual examples


$$\nabla \mathcal{L} = \frac{1}{N} \sum_{i=1}^N \nabla L(\mathbf{y}_i, \hat{\mathbf{y}}_i)$$

Stochastic gradient descent (SGD)

SGD: estimate the gradient using a subset of the examples

- Pick a **single random training example**
- **Estimate** a (noisy) loss on this single training example (the *stochastic* gradient)
- Compute gradient wrt. this loss
- Take a step of gradient descent using the estimated loss

Stochastic gradient descent



Stochastic Gradient Descent (SGD)

SGD Advantages

- Very fast (only need to compute gradient on single example)
- Memory efficient (does not need the full dataset to compute gradient)
- Online (don't need full dataset at each step)

SGD Disadvantages

- Gradient is very noisy, may not always point in correct direction
- Convergence can be slower


In practice: **Mini-batch SGD**

- Estimate gradient on small batch of training examples (say 50)
- Known as **mini-batch stochastic gradient descent**

Vanilla mini-batch SGD

$$\theta_t = \theta_{t-1} - \alpha \underbrace{\nabla_{\theta} \mathcal{L}(\theta_{t-1})}_{\text{Evaluated on a mini-batch}}$$




Momentum

Velocity  $v_t = \gamma v_{t-1} + \alpha \nabla_{\theta} \mathcal{L}(\theta_{t-1})$

$$\theta_t = \theta_{t-1} - v_t$$

2x memory for parameters!

Momentum



$$v_t = \gamma v_{t-1} + \alpha \nabla_{\theta} \mathcal{L}(\theta_{t-1})$$

$$\theta_t = \theta_{t-1} - v_t$$

2x memory for parameters!

Adagrad


Adapts the learning rate for each of the parameters based on sizes of previous updates.

- Scales updates to be larger for parameters that are updated less
- Scales updates to be smaller for parameters that are updated more

Store sum of squares of gradients so far in diagonal of matrix G_t

$$G_t = \sum_{i=0}^t \text{diag}(\nabla \mathcal{L}(\theta)_i)^2$$

Gradient of loss at timestep i



Update rule: $\theta_t = \theta_{t-1} - \alpha G_t^{-\frac{1}{2}} \nabla \mathcal{L}(\theta_{t-1})$


RMSProp

Modification of Adagrad to address aggressively decaying learning rate.

Instead of storing sum of squares of gradient over all time steps so far, use a **decayed moving average** of sum of squares of gradients

$$G_t = \gamma G_{t-1} + (1 - \gamma) \text{diag}(\nabla \mathcal{L}(\theta))^2$$

Update rule: $\theta_t = \theta_{t-1} - \alpha G_t^{-\frac{1}{2}} \nabla \mathcal{L}(\theta_{t-1})$



Adam

Combines momentum and RMSProp

Keep decaying average of both first-order moment of gradient (momentum) and second-order moment (like RMSProp)

First-order: $v_t = \gamma_1 v_{t-1} + (1 - \gamma_1) \nabla \mathcal{L}(\theta_{t-1})$

Second-order: $G_t = \gamma_2 G_{t-1} + (1 - \gamma_2) \text{diag}(\nabla \mathcal{L}(\theta))^2$

Update rule: $\theta_t = \theta_{t-1} - \alpha G_t^{-\frac{1}{2}} v_t$

3x memory!

Summary

We need an algorithm to find **good weight configurations**.

This is an unconstrained continuous **optimization problem**.

We can use standard iterative optimization methods like **gradient descent**.

To use gradient descent, we need a way to find the **gradient of the loss with respect to the parameters** (weights and biases) of the network.

Error **backpropagation** is an efficient algorithm for finding these gradients.

Basically an application of the multivariate **chain rule** and **dynamic programming**.

In practice, computing the full gradient is expensive. Backpropagation is typically used with **stochastic gradient descent**.

Outline

1. RL with Neural Networks
2. Loss functions
3. Backpropagation
4. Optimizers

Undergradese

What undergrads ask vs. what they're REALLY asking

"Is it going to be an open book exam?"

Translation: "I don't have to actually memorize anything, do I?"

"Hmm, what do you mean by that?"

Translation: "What's the answer so we can all go home."

"Are you going to have office hours today?"

Translation: "Can I do my homework in your office?"

"Can i get an extension?"

Translation: "Can you re-arrange your life around mine?"

"Is this going to be on the test?"

Translation: "Tell us what's going to be on the test."

"Is grading going to be curved?"

Translation: "Can I do a mediocre job and still get an A?"

JORGE CHAM © 2008

